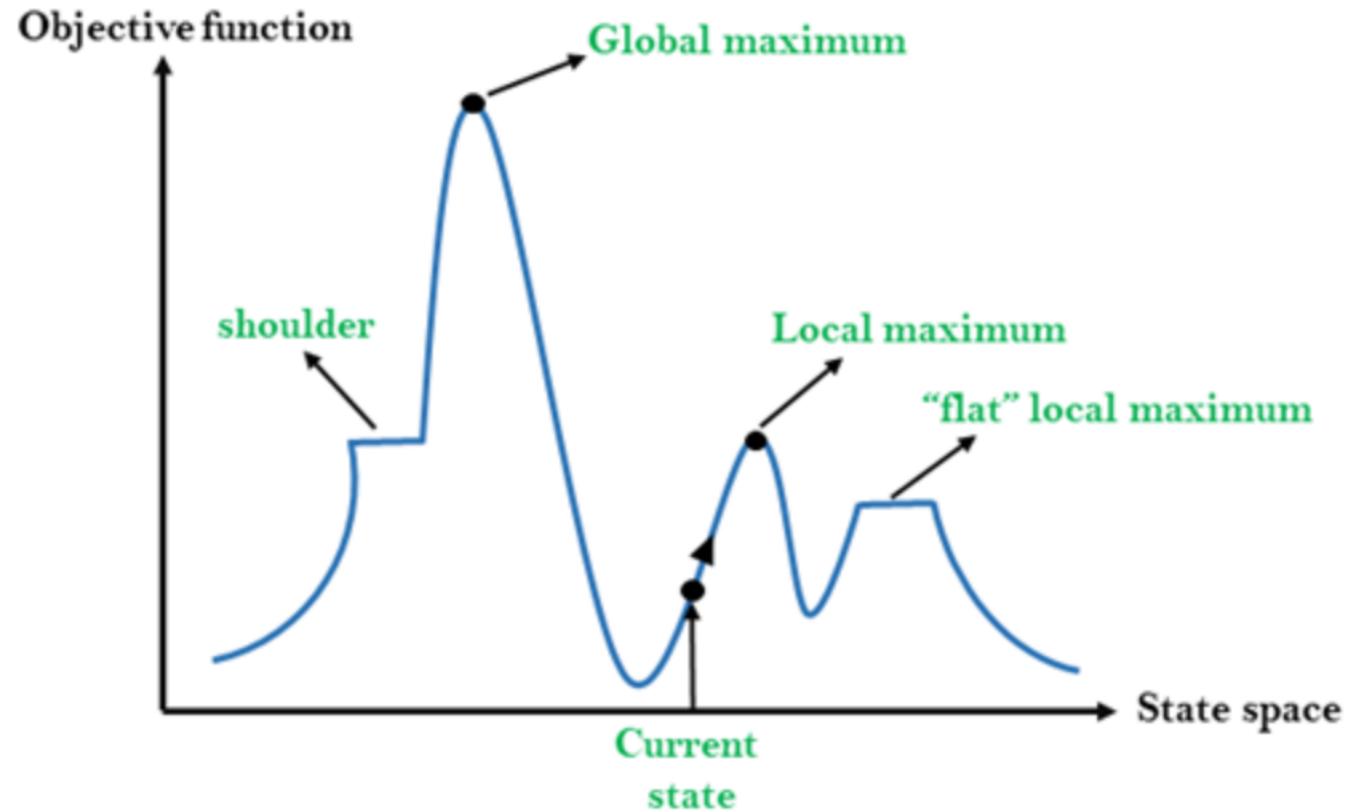


CSCI 3202: Intro to Artificial Intelligence

Lecture 12: Local Search

Rachel Cox
Department of Computer Science

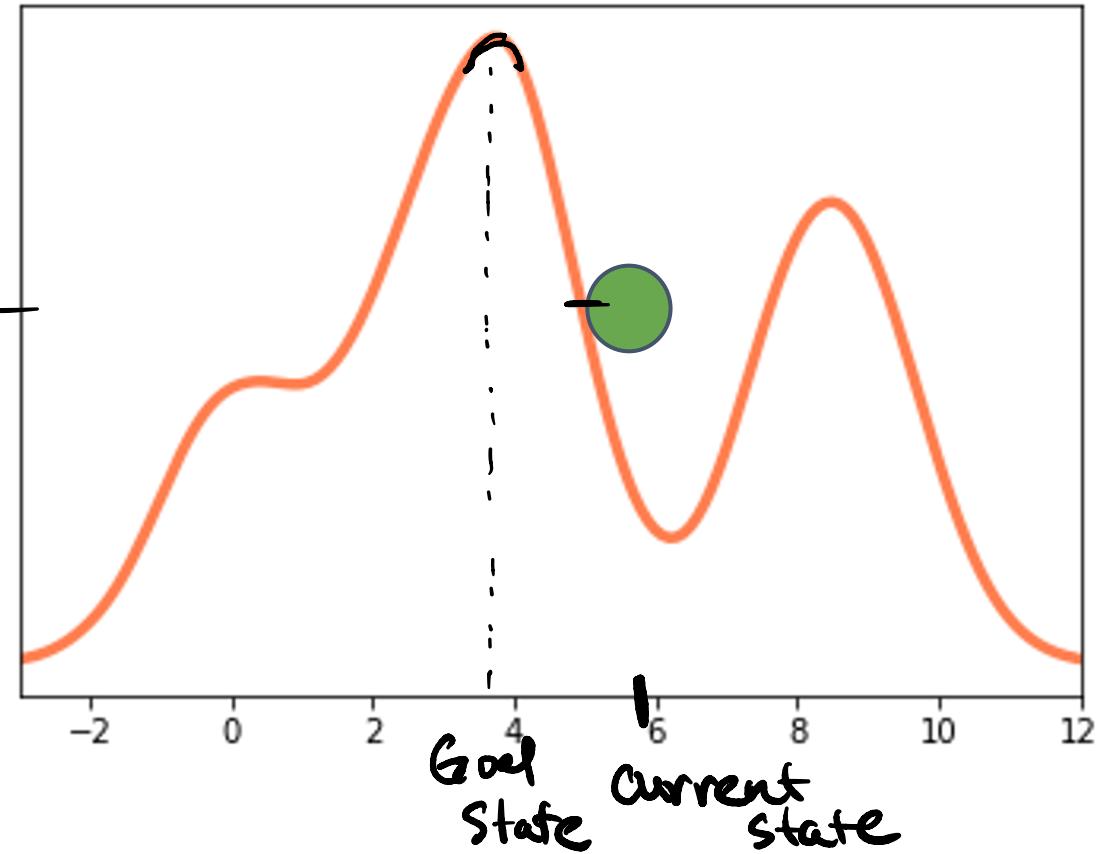


Local Search

Suppose the path to the goal doesn't matter to us.

- look at our neighboring states
- compute the objective function value of the neighboring states
- choose state with max value to move to.

objective
function
value
of the
current
State



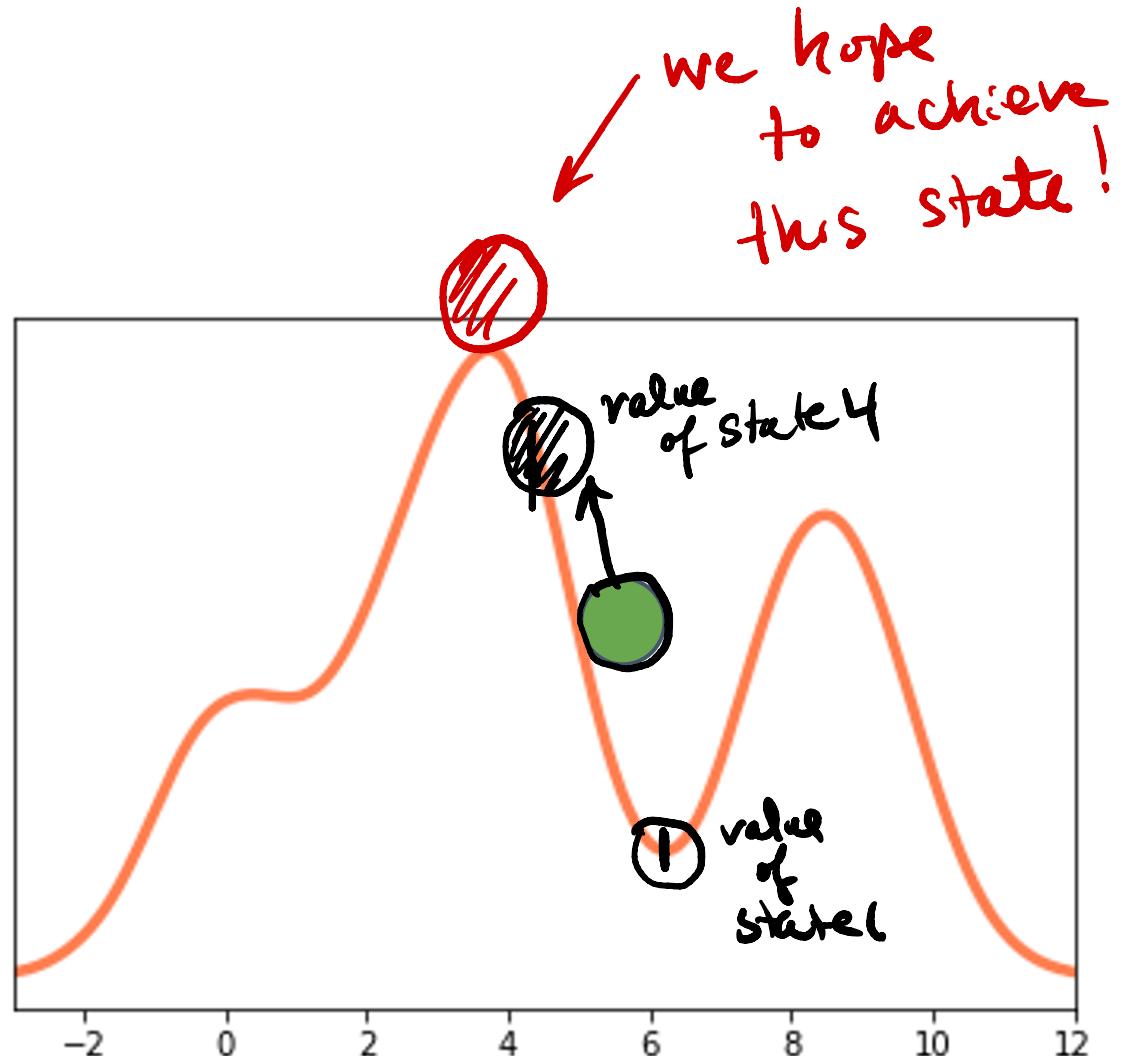
Local Search

Local search:

- Use only the current node/state ←
- And consider only moves to neighbor states
- Super memory efficient (only care about current state)
- Used widely for **optimization**

Find the “best” state according to some **objective function**

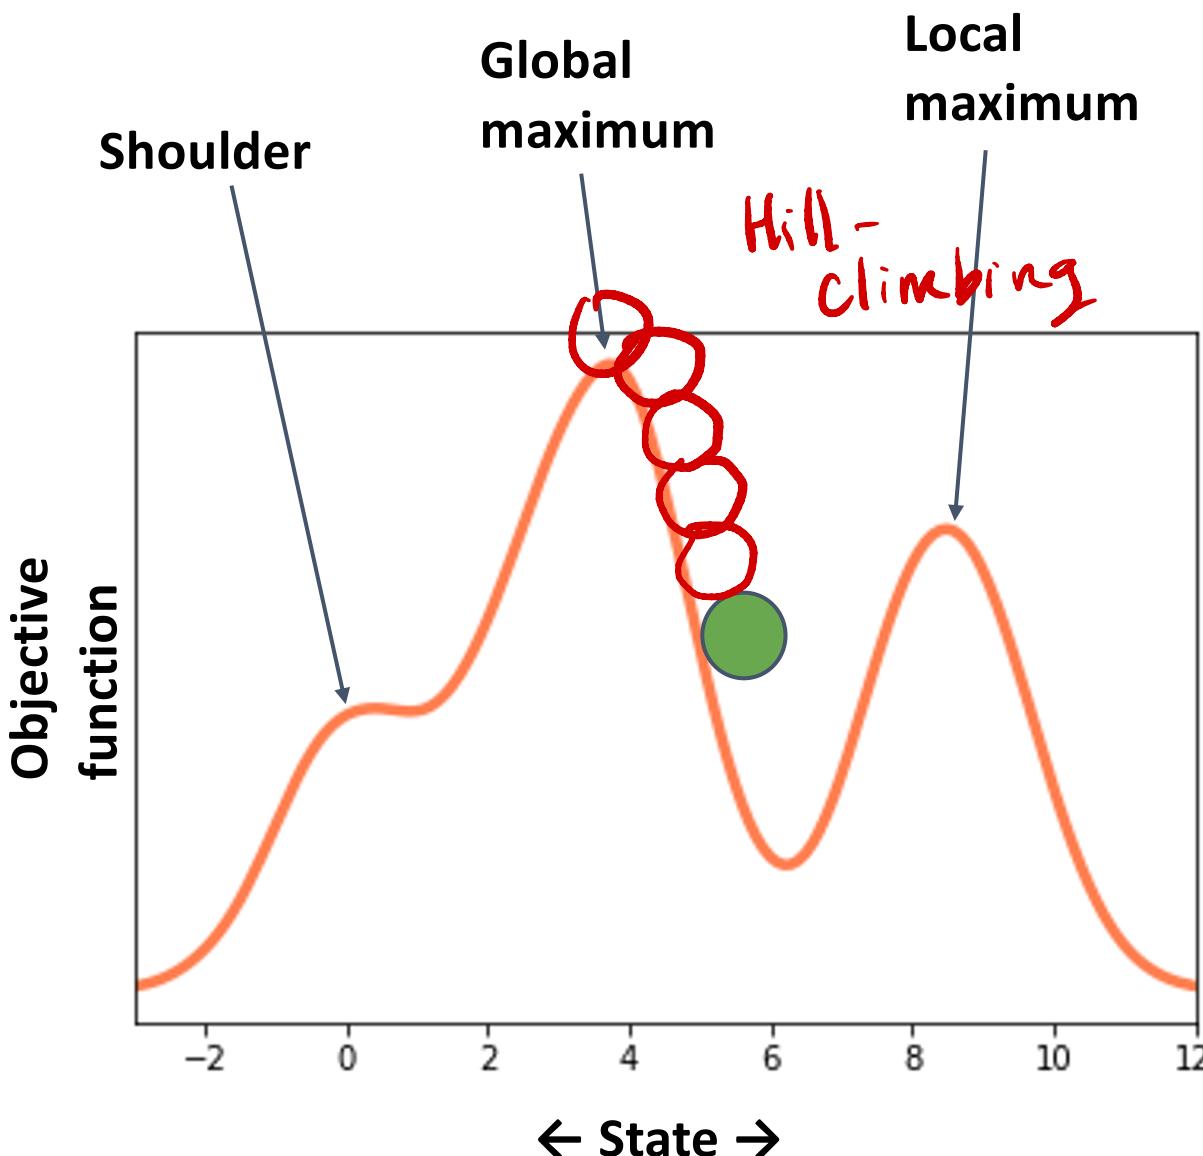
(think: minimizing a loss, or maximizing a utility)



Local Search

Local search:

- Some terminology...
- A complete local search algorithm will always find the goal if one exists.
- An optimal local search algorithm will specifically find the global max or min.

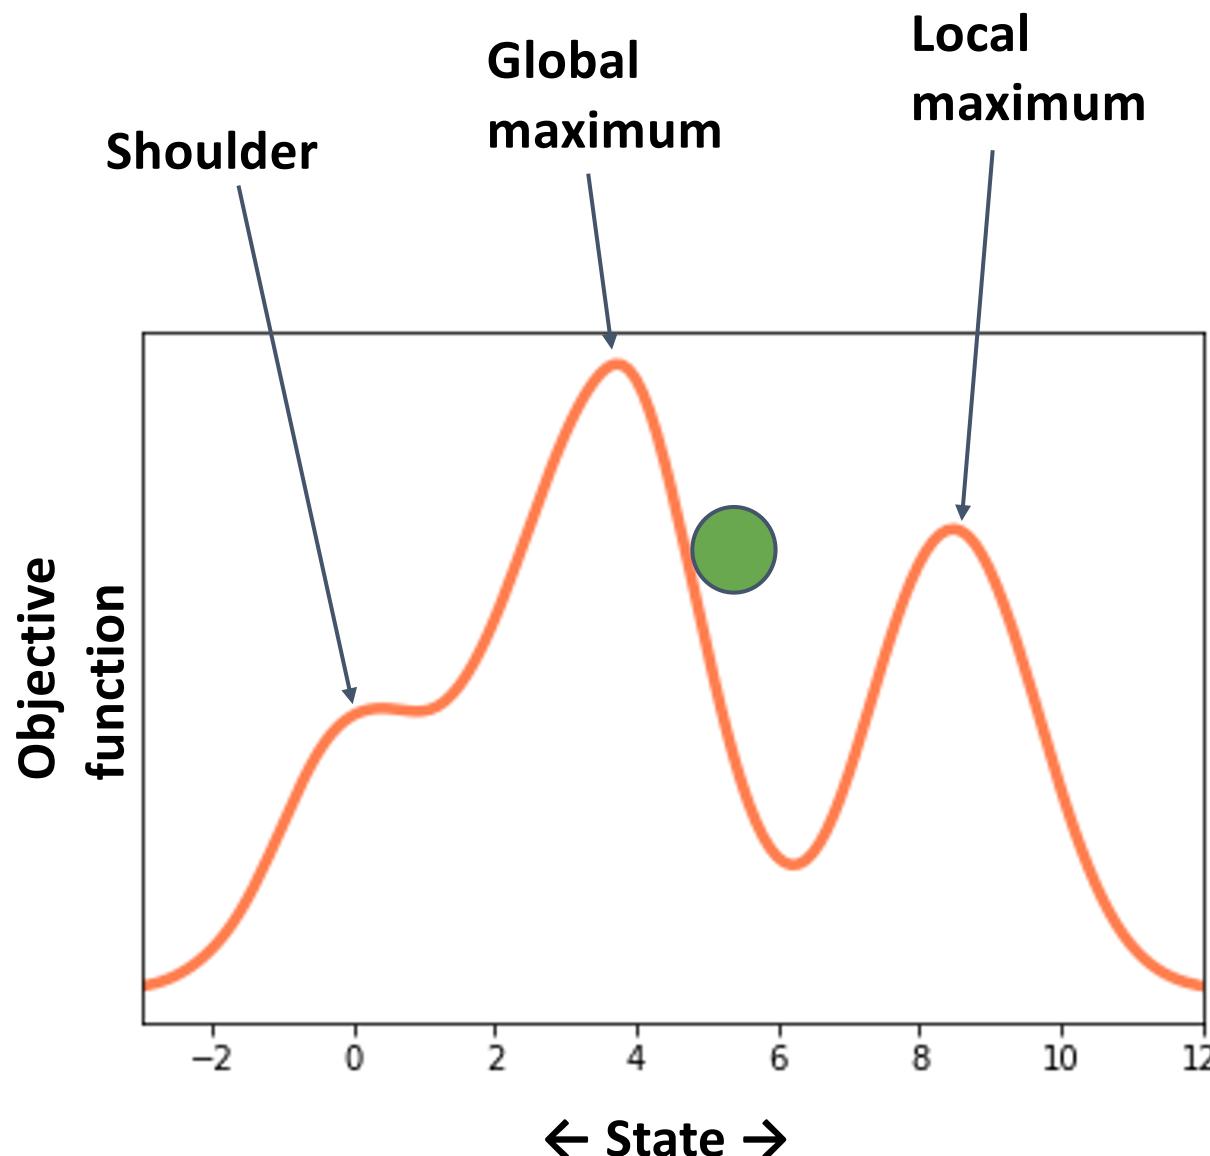


Hill Climbing

AKA: steepest ascent

Some words...

- 1) Start in some initial state
- 2) Get the neighbors of that state
- 3) Move to the neighbor with the best value of the objective function



Hill Climbing

Some potential issues...

- Hit a shoulder?
(or local max.)

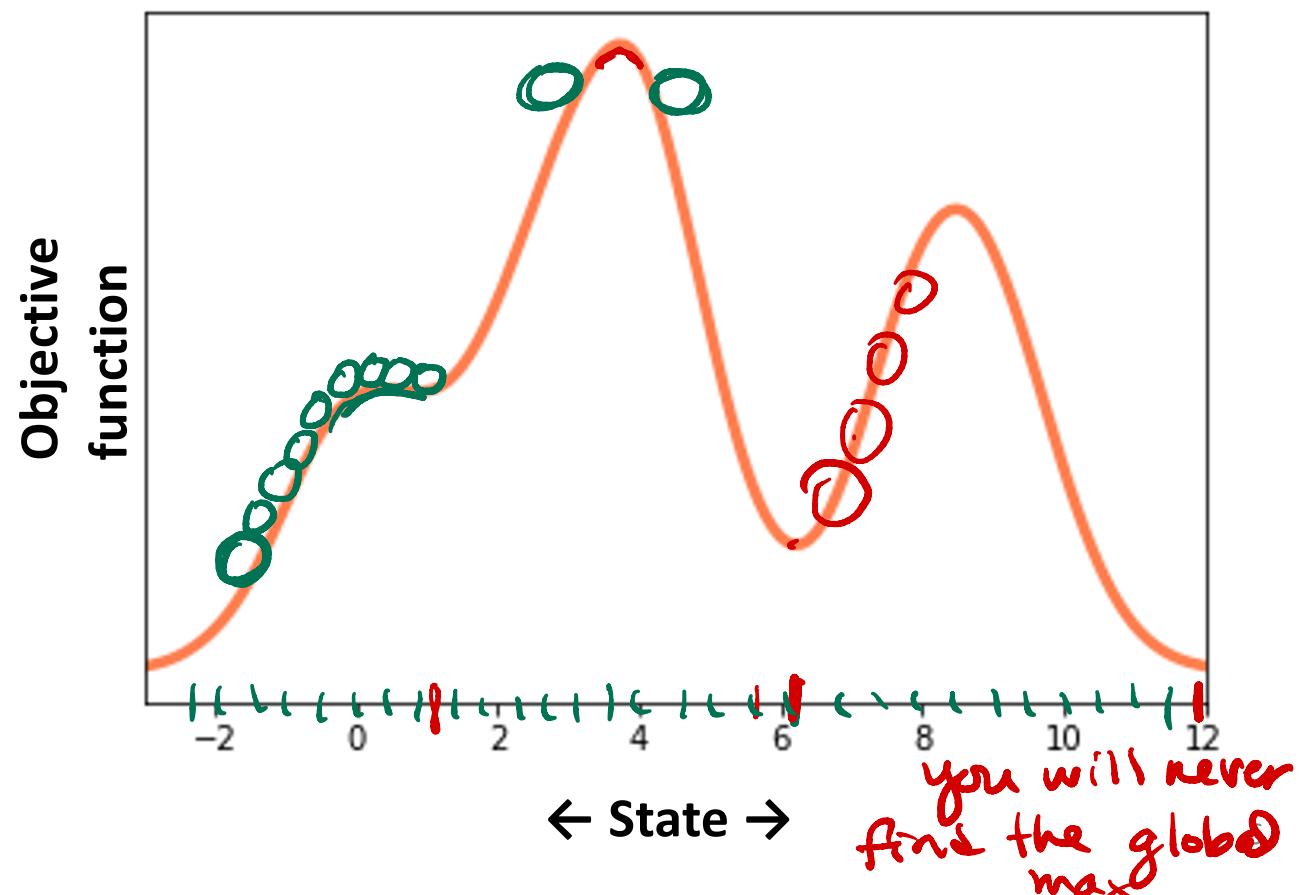
→ Permit sideways moves?

→ “Convergence” criteria?

→ Stochastic hill climbing

→ Random restart

→ simulated annealing



Hill Climbing

Some pseudocodes...

```
class state:  
  
    # we need to be able to keep track of states and their associated  
    # objective function values  
  
class problem:  
  
    # We want to include everything that defines the problem here:  
    # - initial state —  
    # - current state —  
    # - what is the objective function we're trying to maximize/minimize?  
    # - what are the choices of action that we have?  
    # - what do we need to know to select our action? }  
  
def hill_climb(problem, number of iterations):  
  
    for t in some number of iterations:  
        1. get a list of our available moves  
        2. which move optimizes the objective function? —  
        3. do that move; update our state —  
        4. possible goal/convergence check }
```

Annealing – Metallurgy

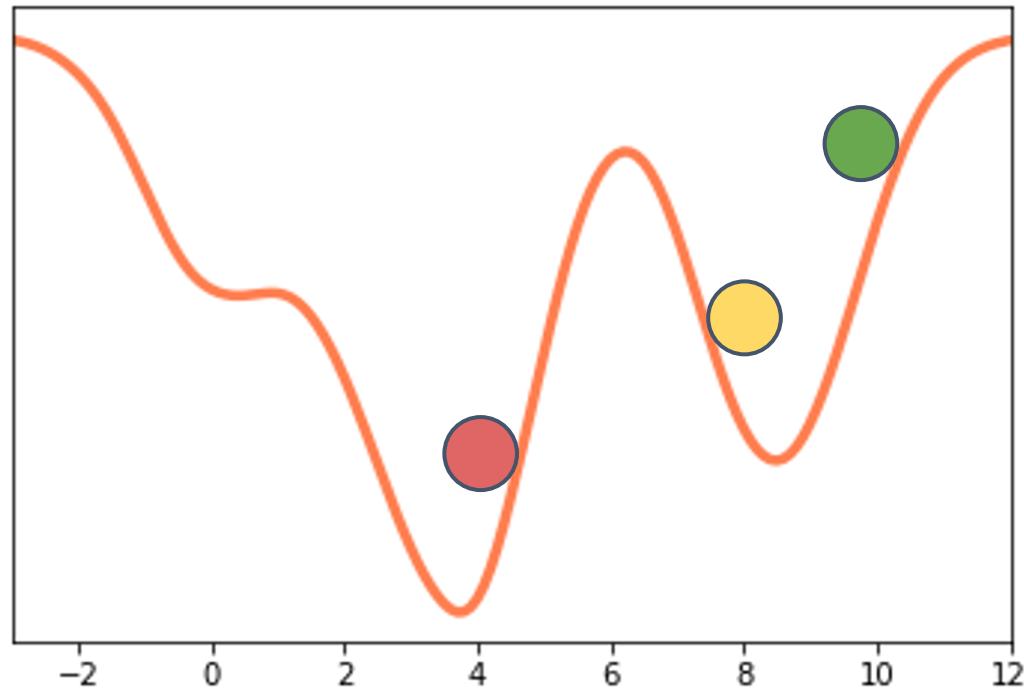
Visualization of what happens when you heat treat (e.g. Anneal) a metal: [Video](#)



Simulated Annealing

Trying to find global minimum (in this case)

- Propose **random** moves
- Accept with a probability that depends on:
 - 1) "temperature"
 - 2) how "good" the move is
 - Accept move without question if it's better than our current situation
 - Accept move with some probability < 1 if it's worse



Simulated Annealing

Visualization of Simulated Annealing: [Video](#)

Meant to mimic the slow cooling of metals.

Each new solution x_j is accepted with a temperature dependent probability.

probability function \rightarrow

$$P_T = \begin{cases} 1 & \text{if } f(x_j) \leq f(x_i) \\ e^{\frac{f(x_i) - f(x_j)}{kT}} & \text{if } f(x_j) \geq f(x_i) \end{cases}$$

T - current Temperature
k - Boltzmann constant

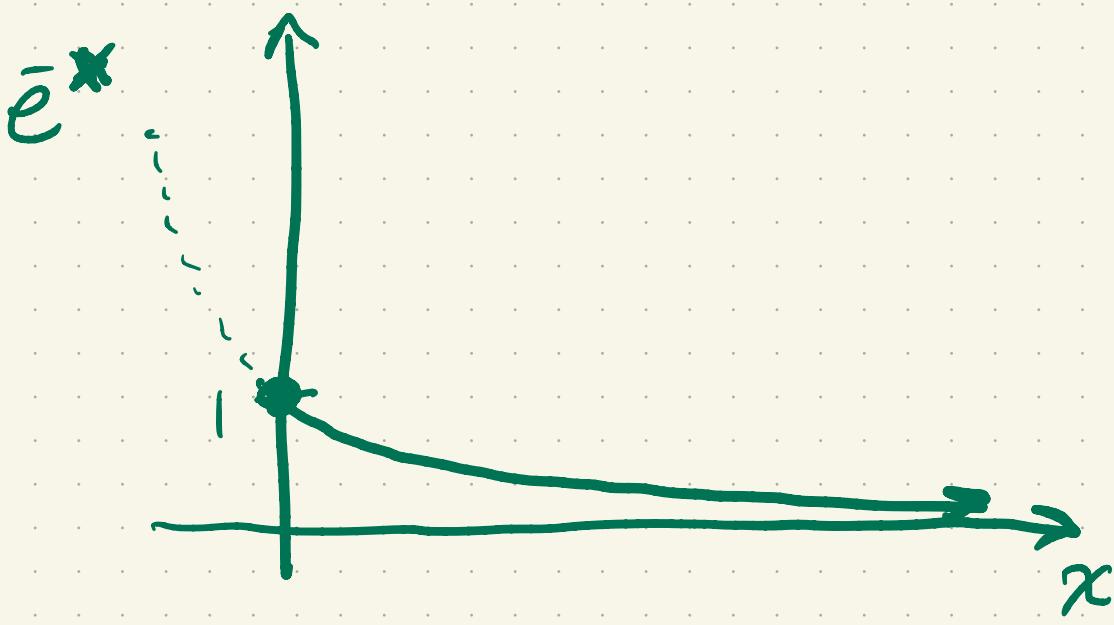
new state is better
→ move with prob 1

f - objective function value

new is worse,
you still might move there
with prob > 0

x_i, x_j states

- ❖ Further reading - optimization is a huge area, we are only scratching the surface



e^{-x} will always provide an output between 0 and 1

s^* : proposed state

s_t : current state

minimize

$$\underline{f(s^*)} > \underline{f(s_t)}$$

$$P_{\text{accept}} = e^{\frac{-\Delta E}{T(t)}}$$

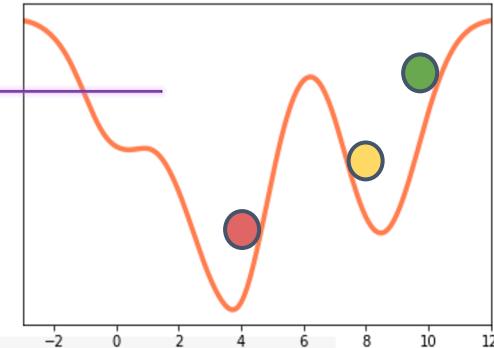
Simulated Annealing

Some pseudocode...

```
def schedule(time):
    '''some sort of mapping from time to temperature, to represent how we should be
    "cooling off" - that is, accepting wacky solutions with lower and lower probability'''
    # math goes here
    return temperature

def simulated_annealing(problem, some number of iterations):

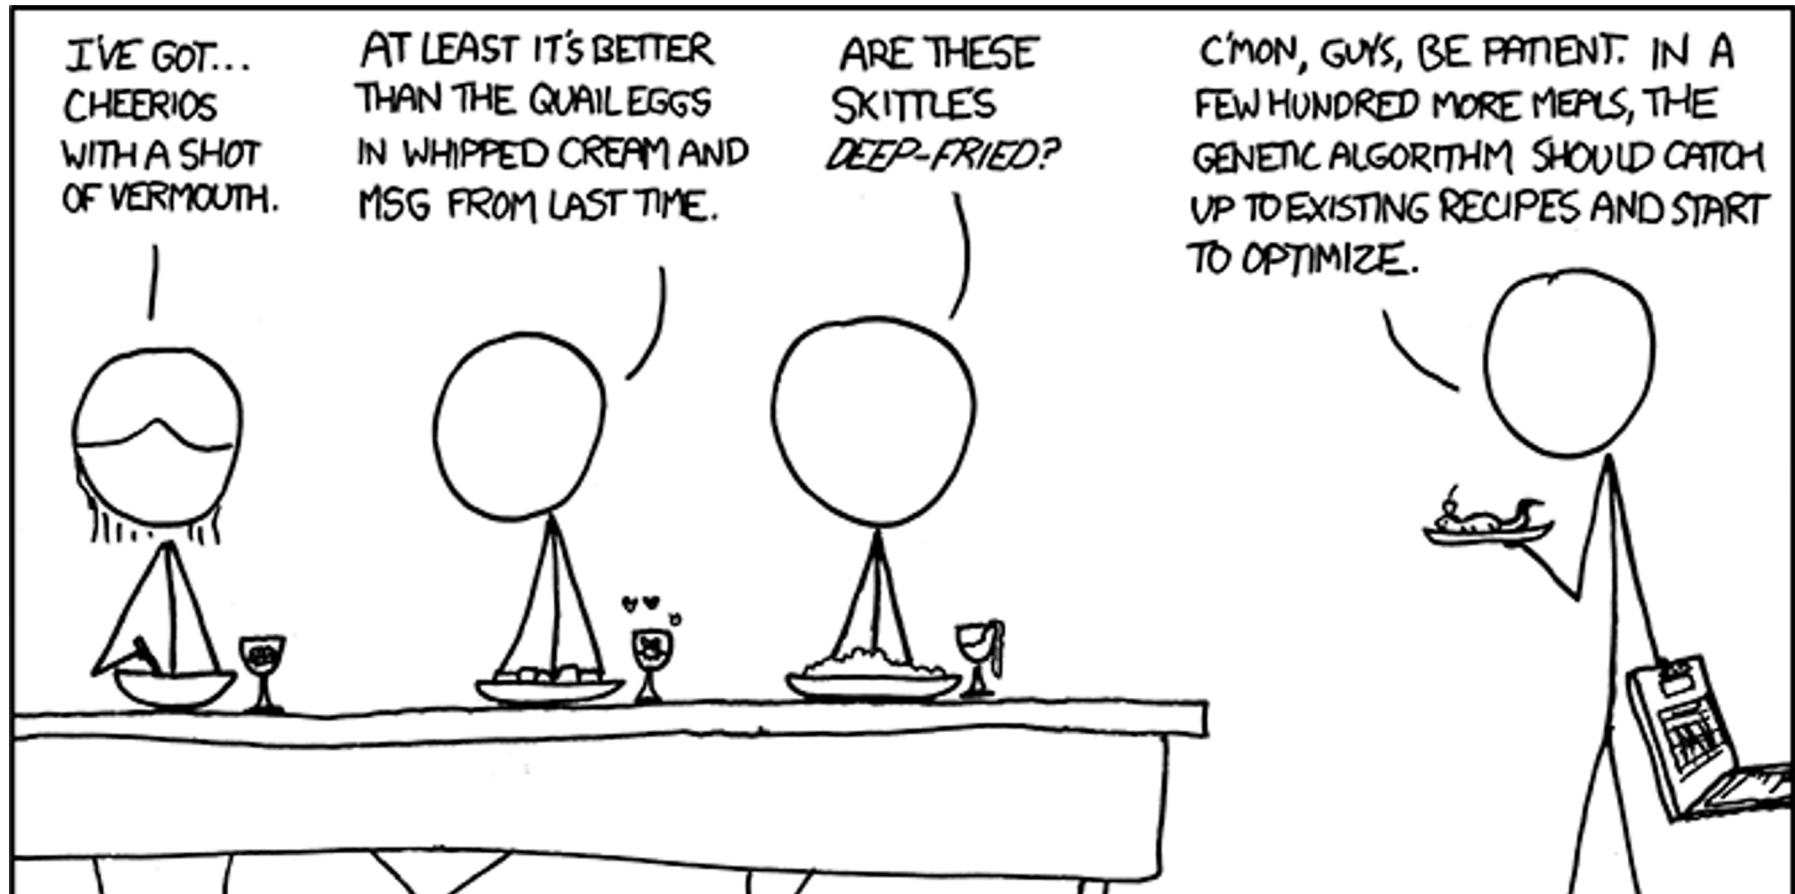
    for t in some number of iterations:
        #1. update the "temperature", T(t) = schedule(time)
        #2. which moves can we make from the current node?
        #3. pick a random move
        #4. calculate difference in objective function between
        #   proposed new state and the current state (deltaE)
        #5. if proposed new state is better than the current state,
        #   then accept the proposed move with probability 1
        #6. otherwise...
        #       ACCEPT the move with probability exp(-deltaE/T(t)),
        #       or REJECT with prob 1-exp(-deltaE/T(t))
```



Next Time

- *Optimization!*
- *Genetic algorithms!*

But first...



WE'VE DECIDED TO DROP THE CS DEPARTMENT FROM OUR WEEKLY DINNER PARTY HOSTING ROTATION.