

CSCI 3202: Intro to Artificial Intelligence

Lecture 17 (Review): Midterm

Rachel Cox
Department of Computer Science



Exam Info

What is an agent?

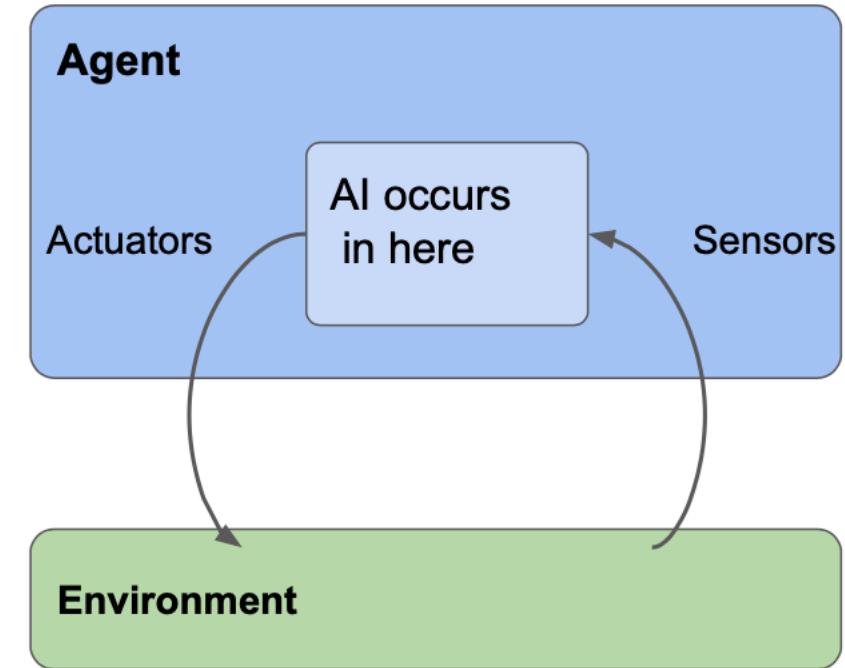
An **agent** is an entity that perceives and acts.

- Perceives via sensors (percepts)
- Acts via actuators (actions)

A **rational agent** is one that does the “right” thing.

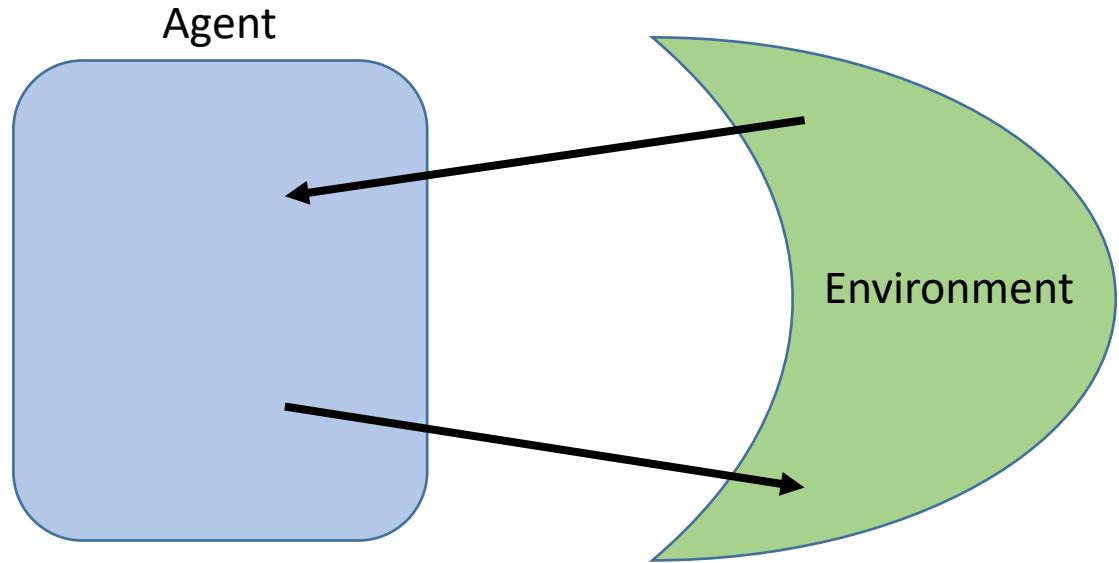
Rationality based on the following four principles:

- 1) performance measure (utility)
- 2) environment familiarity
- 3) actions that are possible
- 4) sequences of percepts (memory)



What is an agent?

Perception-Action Cycle



How does the agent make decisions based on the sensory data?

- ❖ Can't yet build one system that does it all, but we can build a system that does one thing well.

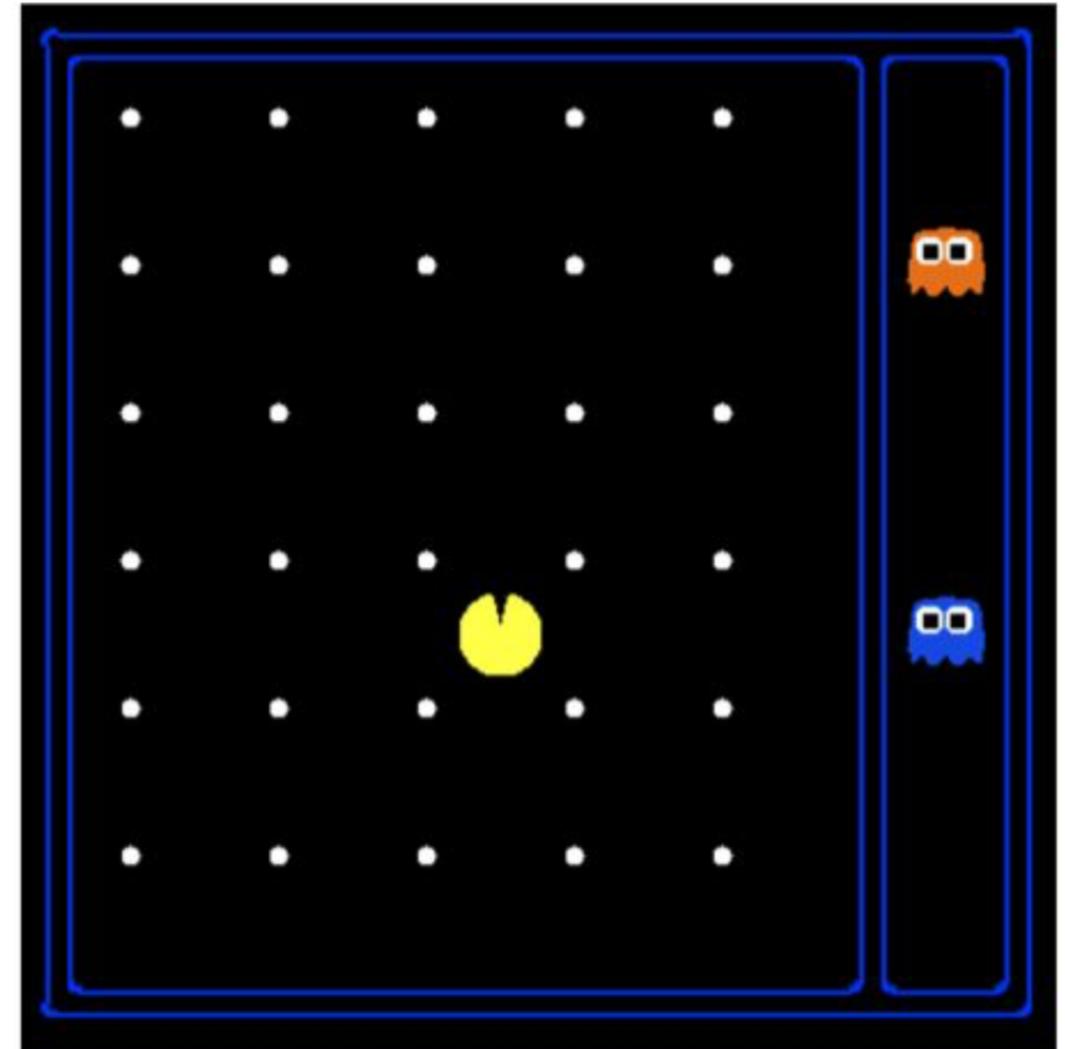
Task Environments

Environment characteristic	Solitaire	Backgammon	Bidding on eBay	Taxi
Fully/partially observable?				- - -
Deterministic/stochastic?		-		
Episodic/sequential task?				
Static / dynamic environment?				
Discrete / continuous?				
Single/multi-agent?				

Search

A **search problem** consists of:

1. State space
2. Transition model
3. Actions
4. Initial state
5. Goal test



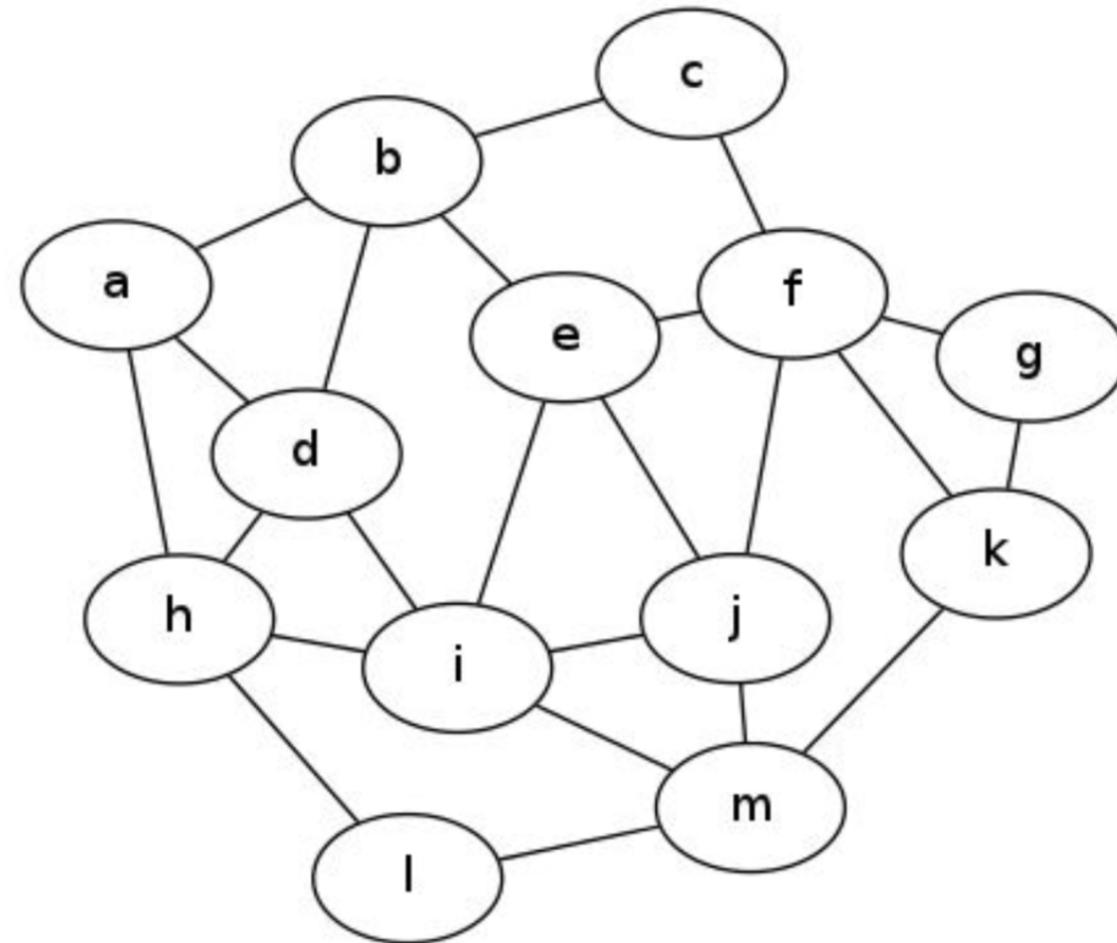
Breadth-First Search

```
function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier  $\leftarrow$  a FIFO queue with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier  $\leftarrow$  INSERT(child, frontier)
```

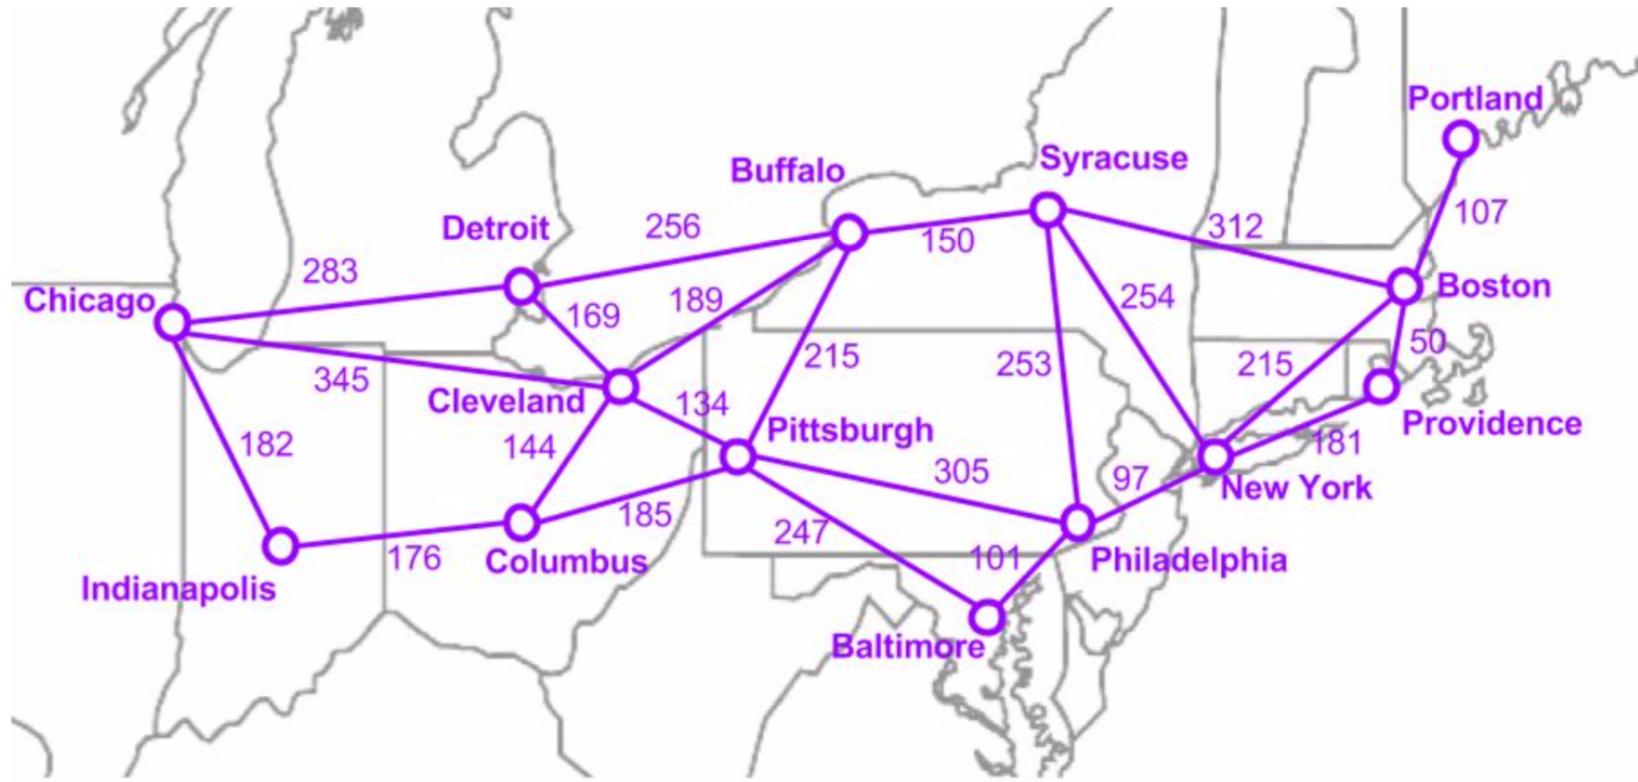
Figure 3.11 Breadth-first search on a graph.

Breadth-first Search (BFS)

Example: Number the nodes in the search graph according to the order in which they would be expanded using BFS to find a path from *a* to *k*. Assume that nodes within a layer are expanded in alphabetical order.



Breadth-first Search (BFS)



Complete?

Optimal?

Depth-First Search

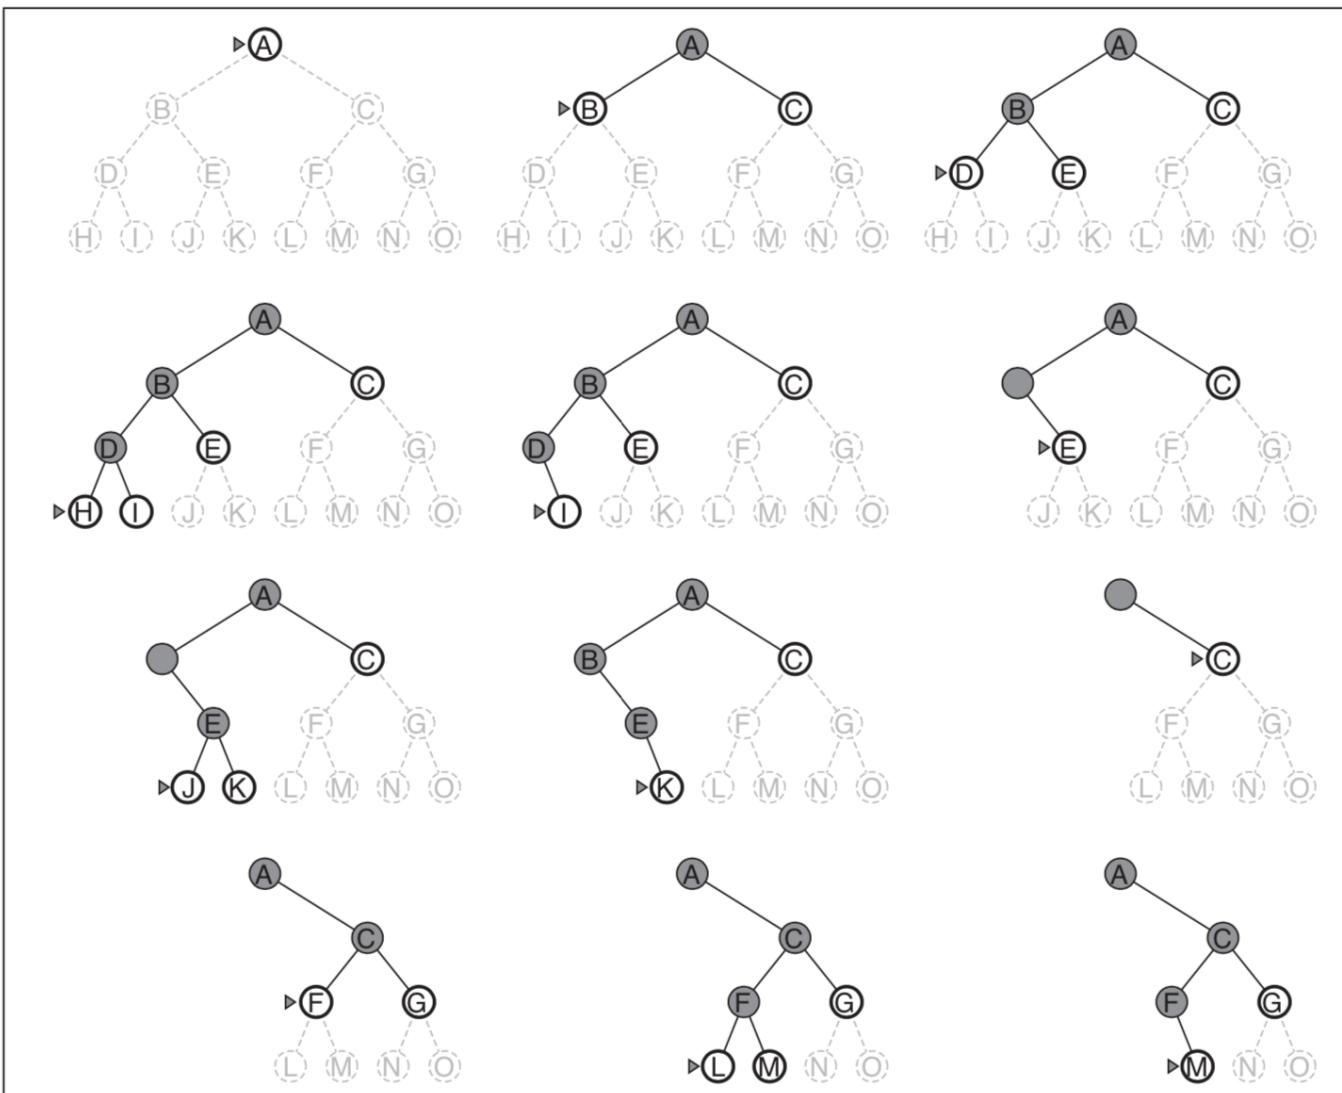
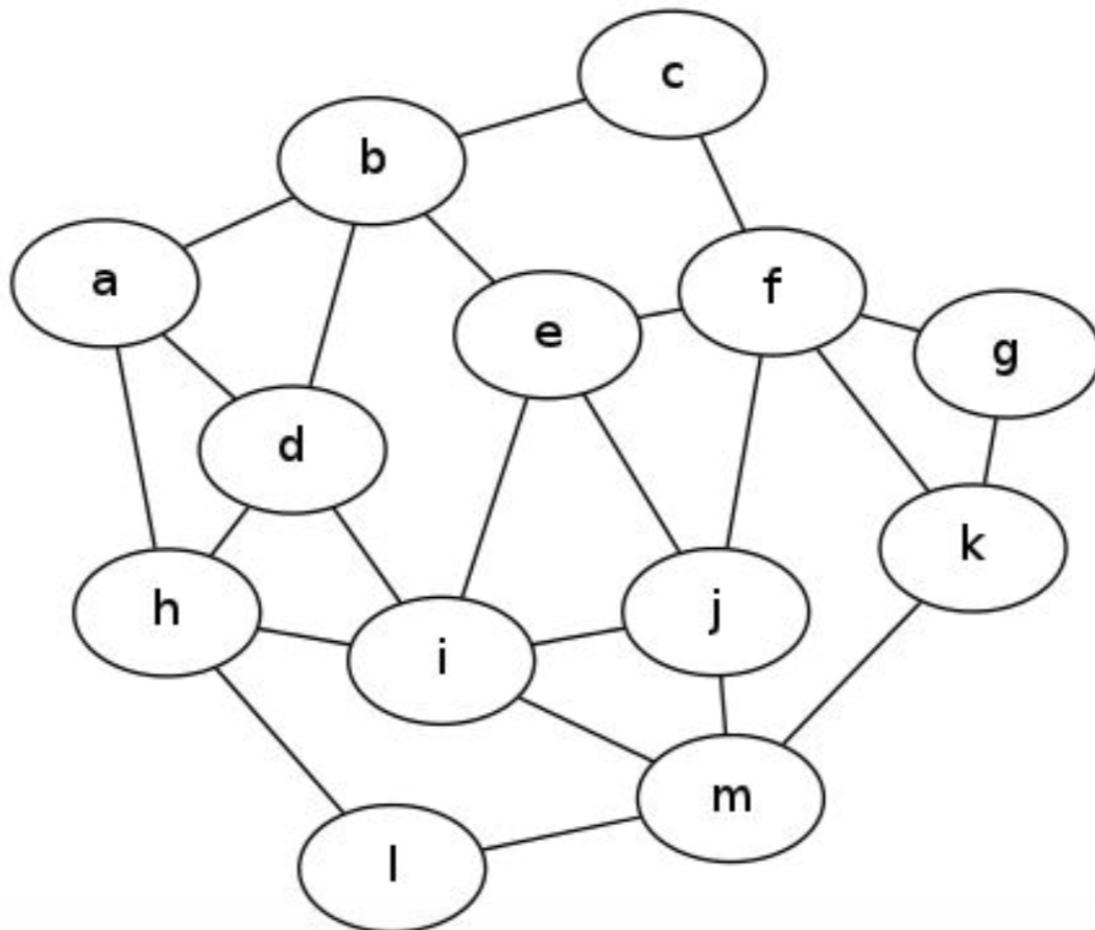


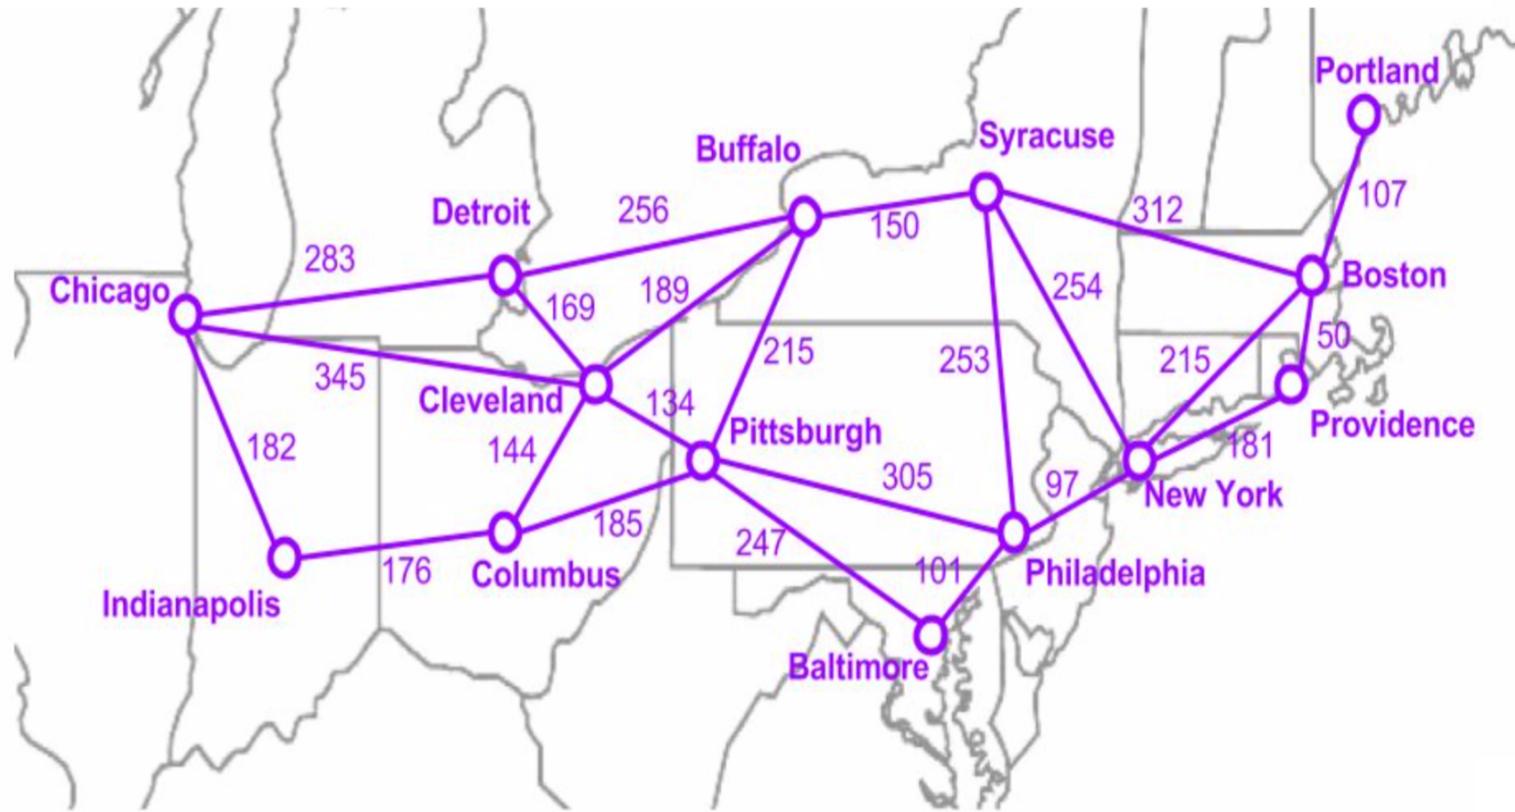
Figure 3.16 Depth-first search on a binary tree. The unexplored region is shown in light gray. Explored nodes with no descendants in the frontier are removed from memory. Nodes at depth 3 have no successors and M is the only goal node.

Depth-First Search (DFS)

Example: Number the nodes in the search graph according to the order in which they would be expanded using DFS to find a path from *a* to *k*. Assume that nodes within a layer are added to the queue by alphabetical order. What is the route that DFS yields, if any?



Depth-First Search (DFS)



Complete?

Optimal?

Uniform Cost Search (UCS)

```
function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
    node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
    frontier  $\leftarrow$  a priority queue ordered by PATH-COST, with node as the only element
    explored  $\leftarrow$  an empty set
    loop do
        if EMPTY?(frontier) then return failure
        node  $\leftarrow$  POP(frontier) /* chooses the lowest-cost node in frontier */
        if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
        add node.STATE to explored
        for each action in problem.ACTIONS(node.STATE) do
            child  $\leftarrow$  CHILD-NODE(problem, node, action)
            if child.STATE is not in explored or frontier then
                frontier  $\leftarrow$  INSERT(child, frontier)
            else if child.STATE is in frontier with higher PATH-COST then
                replace that frontier node with child
```

Figure 3.14 Uniform-cost search on a graph. The algorithm is identical to the general graph search algorithm in Figure 3.7, except for the use of a priority queue and the addition of an extra check in case a shorter path to a frontier state is discovered. The data structure for *frontier* needs to support efficient membership testing, so it should combine the capabilities of a priority queue and a hash table.

A* Search

Uniform-cost search:

$$f(n) = g(n) \quad (\text{cost to get to } n)$$

Greedy:

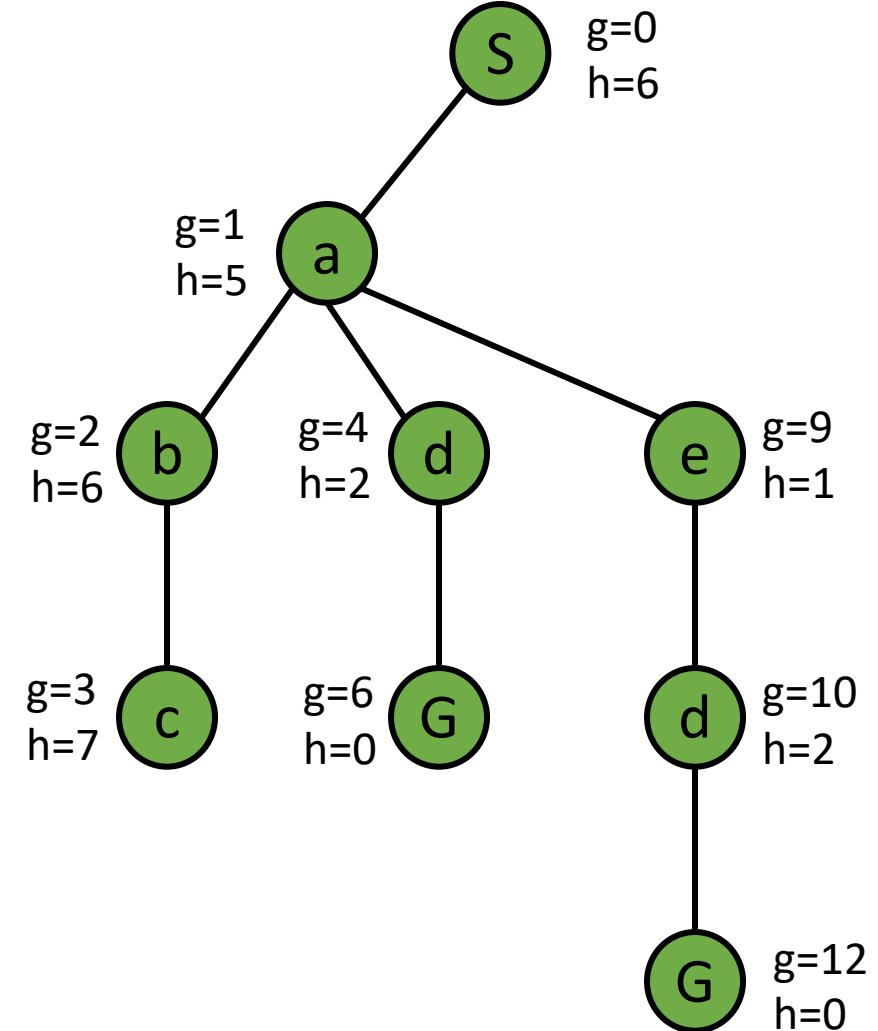
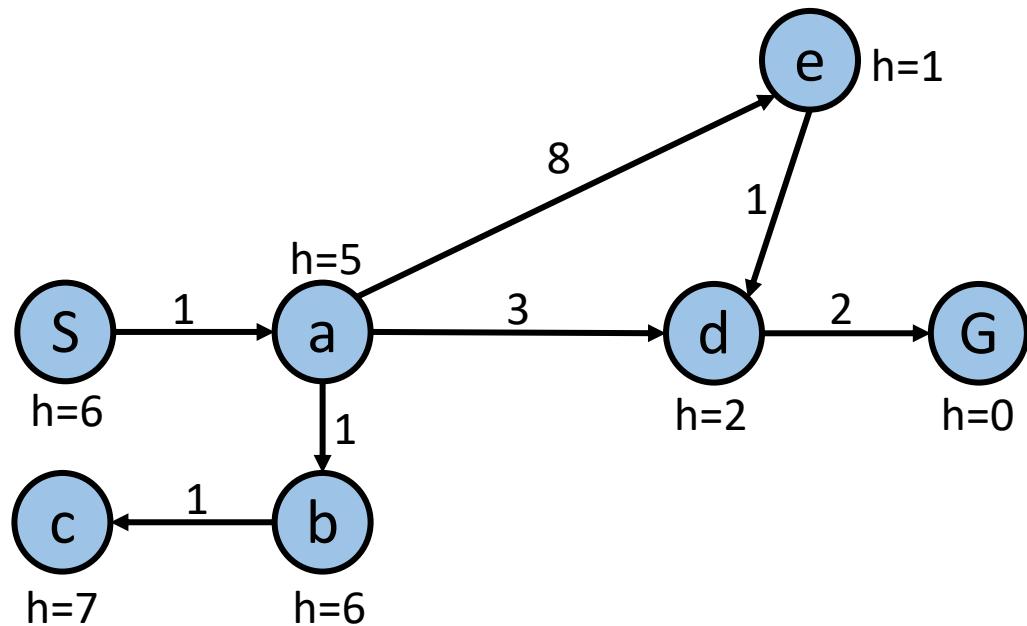
$$f(n) = h(n) \quad (\text{estimated cost to get from } n \text{ to goal})$$

A*:

$$f(n) = g(n) + h(n) \quad (\text{estimated total cost of cheapest solution through } n)$$

A* Search

Example: Compare Uniform Cost, Greedy Search, and A* on the graph below.



A* Search

Consistent: for every node n and successor n' of n , generated by some action a , the estimated cost of reaching the goal from n is no greater than the step cost from n to n' , plus the estimated cost of reaching the goal from n'

- That is: $h(n) \leq c(n, a, n') + h(n')$
- General **triangle inequality** between n , n' , and the goal

A heuristic h is **admissible** (optimistic) if $0 \leq h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to the nearest goal.

A* Search

Any consistent heuristic is also admissible (but not the other way around).

Example: Prove the above statement by induction.

A* Search

Optimality

Conditions for Optimality: Admissibility & Consistency

- $h(n)$ must be **admissible** - an admissible heuristic is one that never overestimates the cost to reach the goal.
- $h(n)$ is **consistent** if, for every node n and every successor n' of n generated by any action a , the estimated cost of reaching the goal from n is no greater than the step cost of getting to n' plus the estimated cost of reaching the goal from n' :

$$h(n) \leq c(n, a, n') + h(n')$$

Optimality of A* Search

A* (graph) is optimal if the heuristic $h(n)$ is consistent.

Based on two key facts:

1. If $h(n)$ is consistent, then the values of $f(n)$ along any path are nondecreasing.
2. Whenever A* selects a node n for expansion, the optimal path to that node has been found.

➤ **So the first goal node to be expanded took the lowest-cost path, and all later goal node expansions are at least as expensive.**

Heuristics

How do we come up with heuristics?

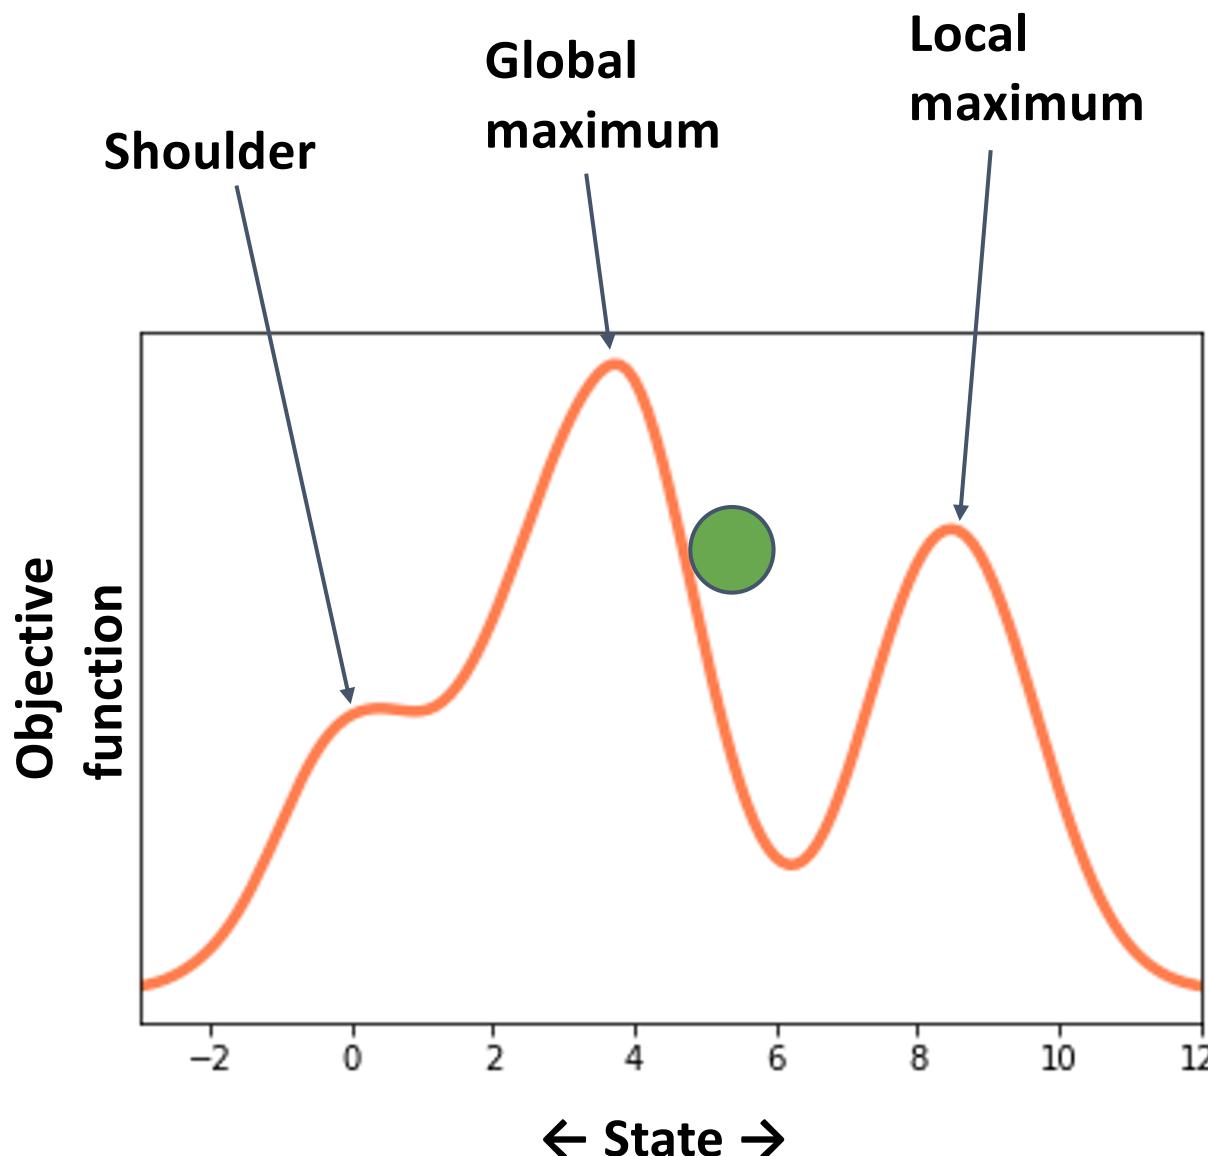
- 1) Generate heuristics from relaxed problems.
- 2) Generate heuristics from sub-problems.
- 3) Learning heuristics from experience.

Hill Climbing

AKA: steepest ascent

Some words...

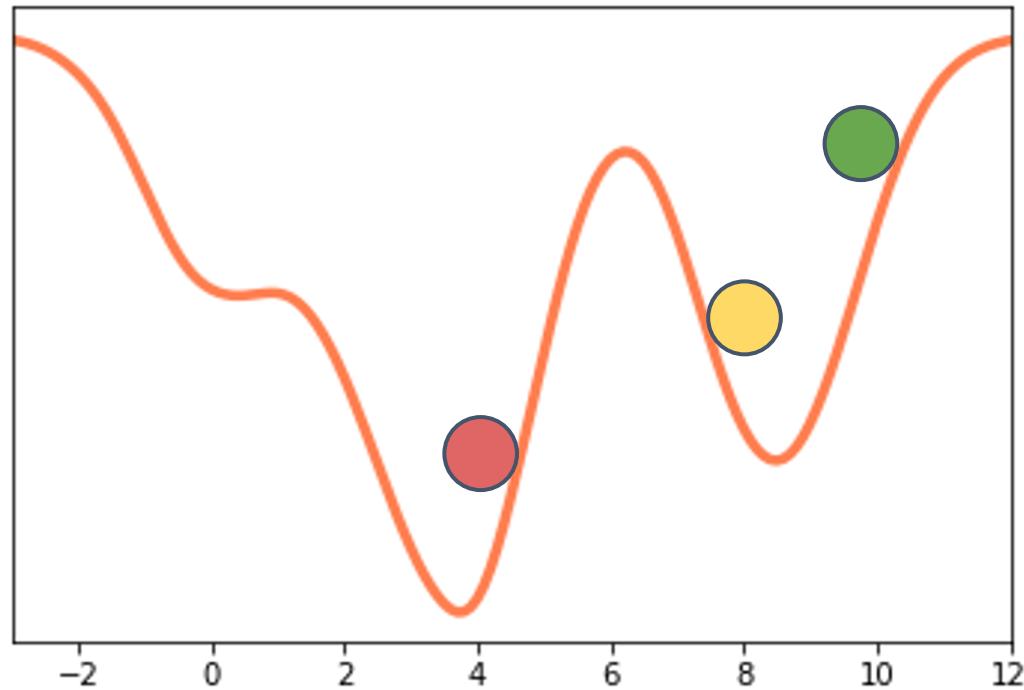
- 1) Start in some initial state
- 2) Get the neighbors of that state
- 3) Move to the neighbor with the best value of the objective function



Simulated Annealing

Trying to find global minimum (in this case)

- Propose **random** moves
- Accept with a probability that depends on:
 - 1) "temperature"
 - 2) how "good" the move is
 - Accept move without question if it's better than our current situation
 - Accept move with some probability < 1 if it's worse



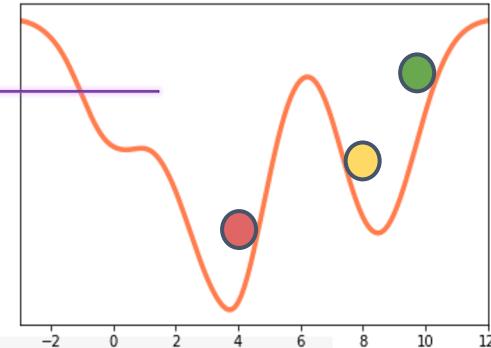
Simulated Annealing

Some pseudocode...

```
def schedule(time):
    '''some sort of mapping from time to temperature, to represent how we should be
    "cooling off" - that is, accepting wacky solutions with lower and lower probability'''
    # math goes here
    return temperature

def simulated_annealing(problem, some number of iterations):

    for t in some number of iterations:
        #1. update the "temperature", T(t) = schedule(time)
        #2. which moves can we make from the current node?
        #3. pick a random move
        #4. calculate difference in objective function between
        #   proposed new state and the current state (deltaE)
        #5. if proposed new state is better than the current state,
        #   then accept the proposed move with probability 1
        #6. otherwise...
        #       ACCEPT the move with probability exp(-deltaE/T(t)),
        #       or REJECT with prob 1-exp(-deltaE/T(t))
```



Random Restart Hill-Climbing

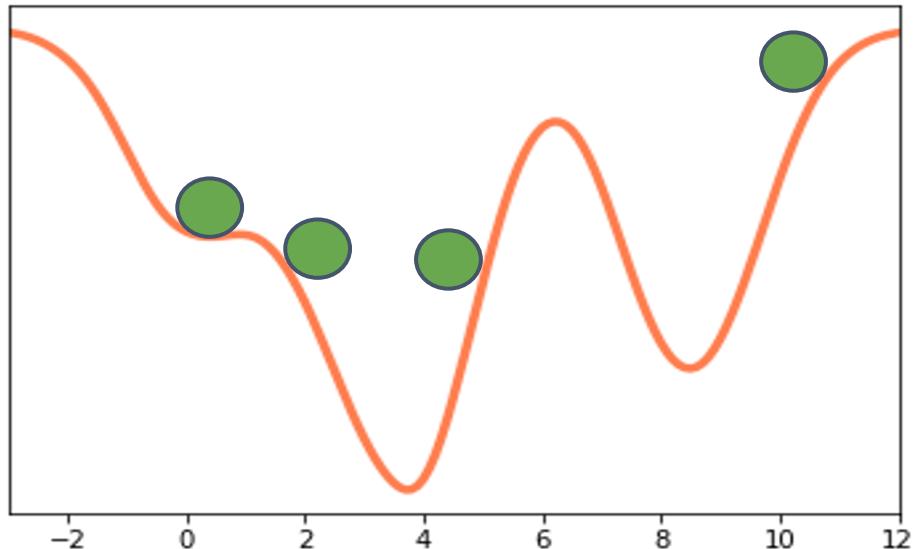
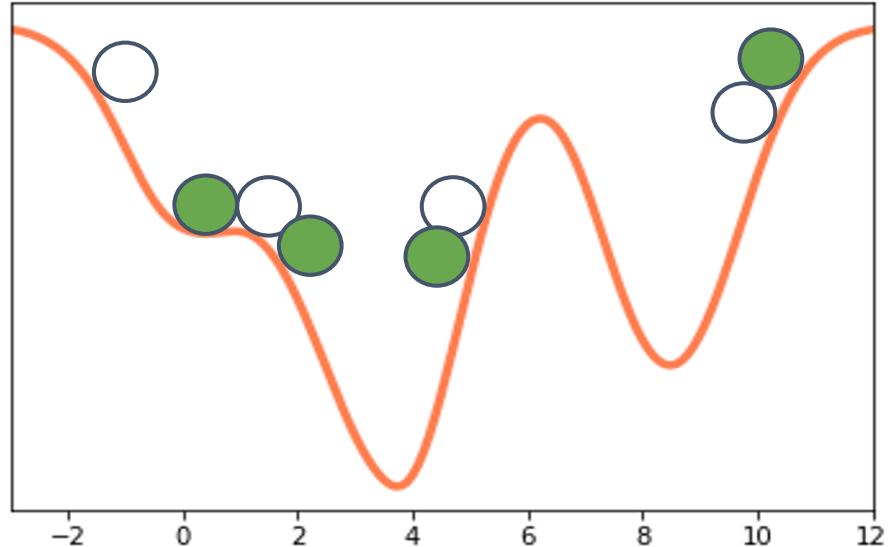
- Conduct a series of hill-climbing searches with a randomized starting point.
- If your random restart initial points end up all being close together, this method may not be any more effective than basic hill-climbing.
- You want to cover the space as much as you can.
- Cost is a constant multiple of hill-climbing - not much more expensive.

Local Beam Search

Fire off k states to explore

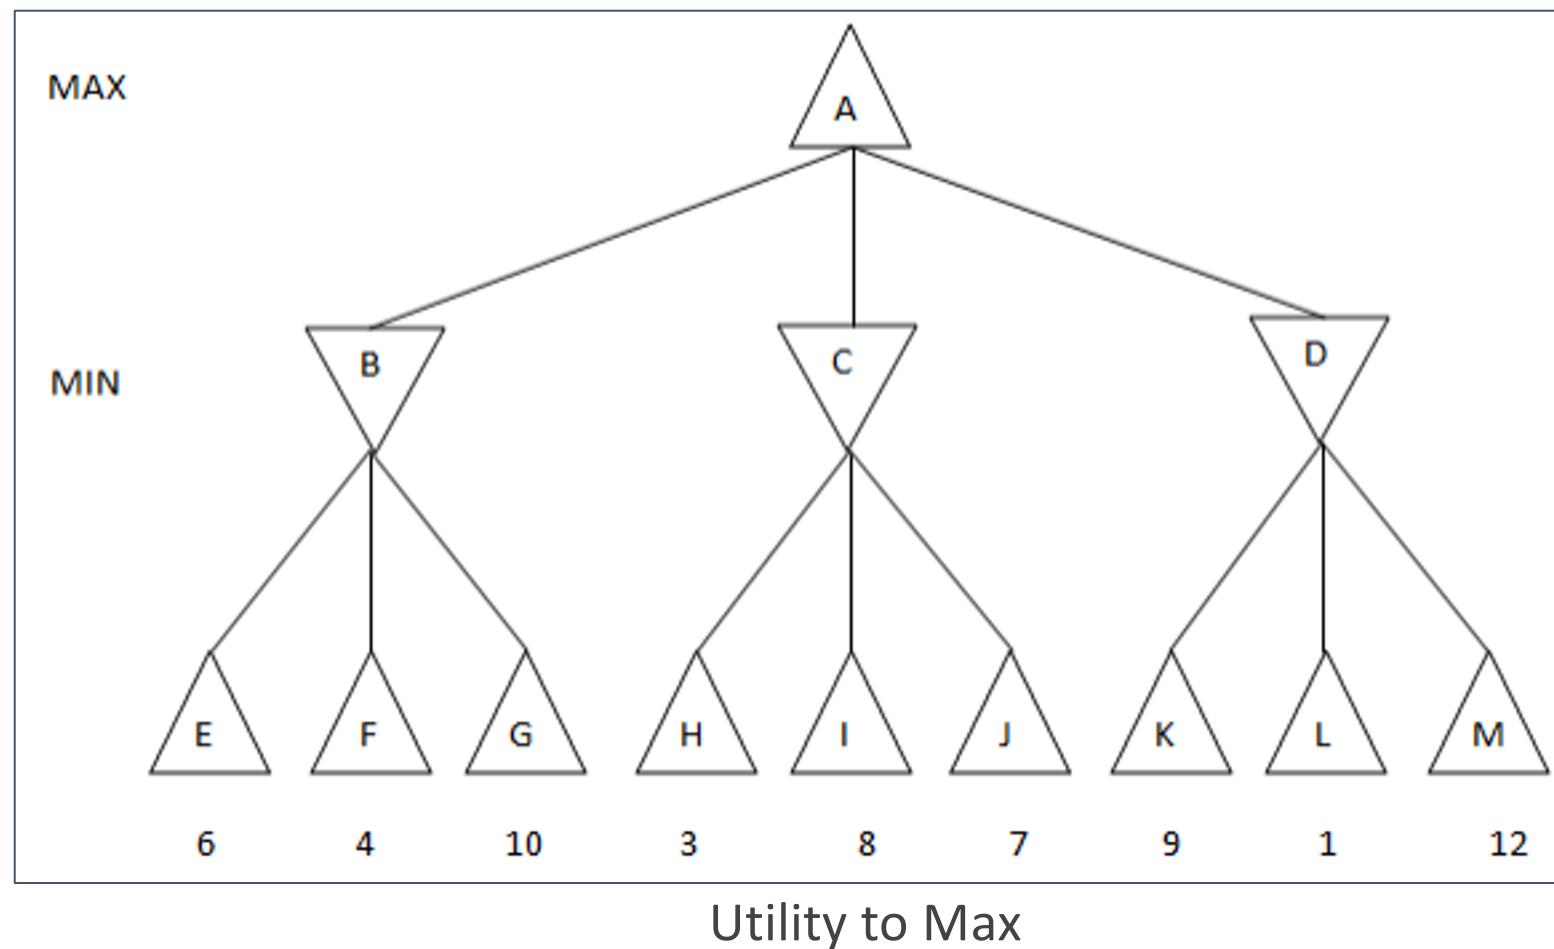
Each step:

- Propose ***all*** successors to each of the k states
- If any of successors is goal -> finished
- Pick k of them to keep exploring in the next step
 - Can be **best k** , or **random** (possibly with some performance-weighting)
- Think of these as **generations** of states...



Minimax

- Two players: **Max** and **Min**
- Alternate making moves
- Optimal strategy from state s determined by **minimax value** of that state.



Minimax

```
def minimax_decision(state):
    all_actions = what are the available actions?
    best_action = action that maximizes min_value(result(action, state))
    return best_action

def min_value(state):
    if terminal_state(state):
        return utility(state)
    value = infinity
    for action in all available actions:
        value = min(value, max_value(result(action, state)))
    return value

def max_value(state):
    if terminal_state(state):
        return utility(state)
    value = -infinity
    for action in all available actions:
        value = max(value, min_value(result(action, state)))
    return value
```

Alpha–Beta Pruning

Doing a full search on the game tree is generally intractable.

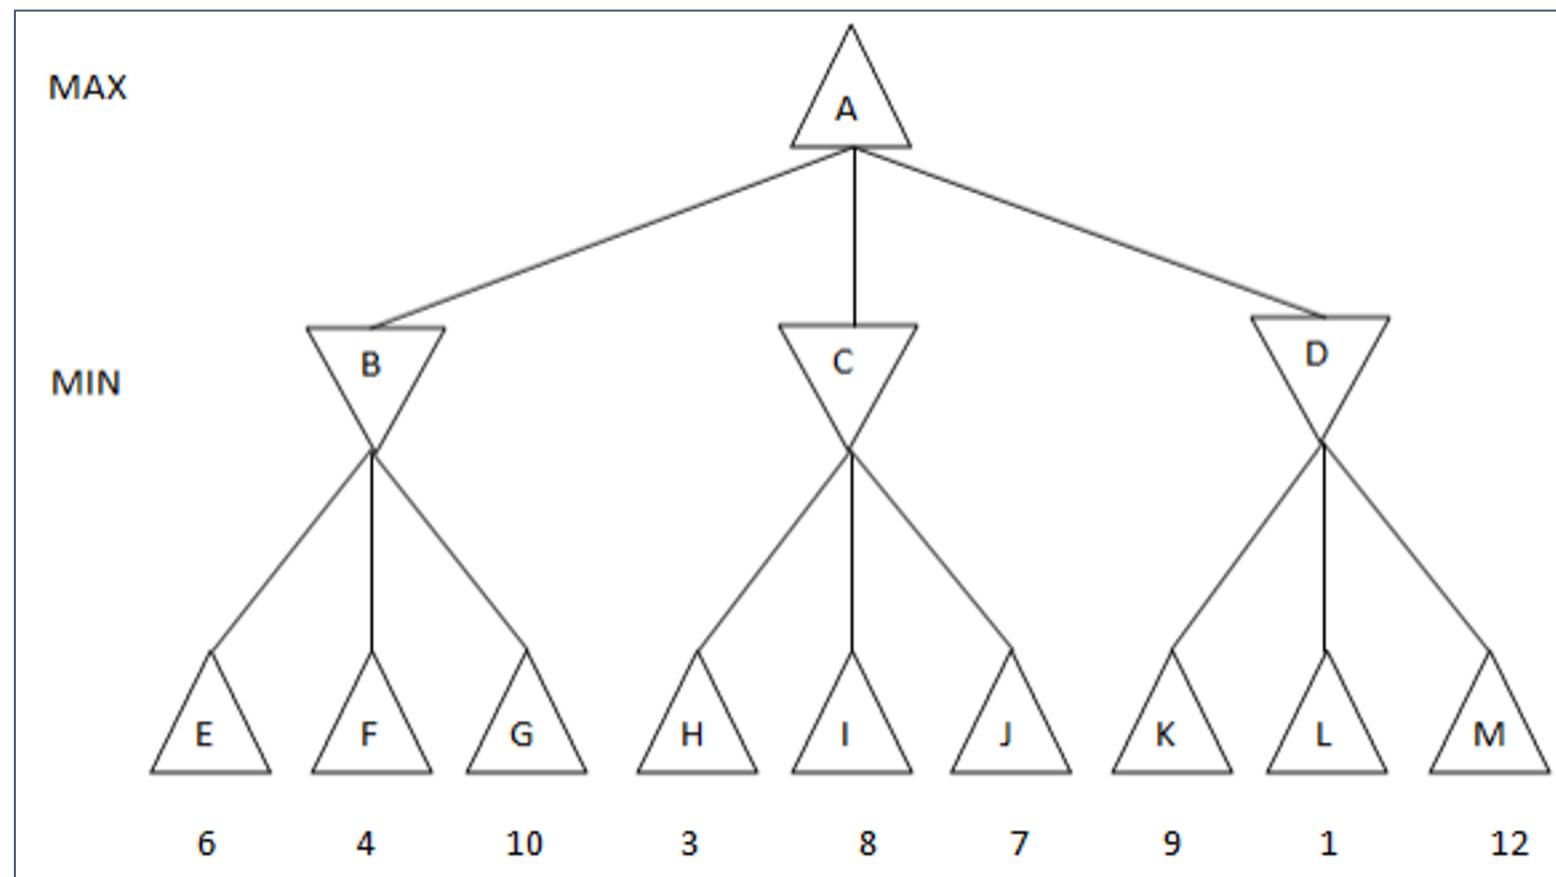
DFS time complexity: $O(b^m)$

m = maximum depth

b = branching factor

BUT if we assume optimal play, then there are some branches of the game tree that we don't need to explore

→ *pruning* those branches



Alpha-Beta Pruning

```
def alphabeta_search(state):
    alpha = -infinity
    beta = +infinity
    value = max_value(state, alpha, beta)
    best_action = action that has utility=value to Max
    return

def max_value(state, alpha, beta):
    if terminal_state(state):
        return utility(state)
    value = -infinity
    for action in all available actions:
        value = max(value, min_value(result(action, state), alpha, beta))
        if value >= beta: return value
        alpha = max(value, alpha)
    return value

def min_value(state, alpha, beta):
    if terminal_state(state):
        return utility(state)
    value = +infinity
    for action in all available actions:
        value = min(value, max_value(result(action, state), alpha, beta))
        if value <= alpha: return value
        beta = min(value, beta)
    return value
```

Alpha–Beta Pruning

Example: Consider the Game Tree to below. Find the resulting value at the root node by following the minimax algorithm.

